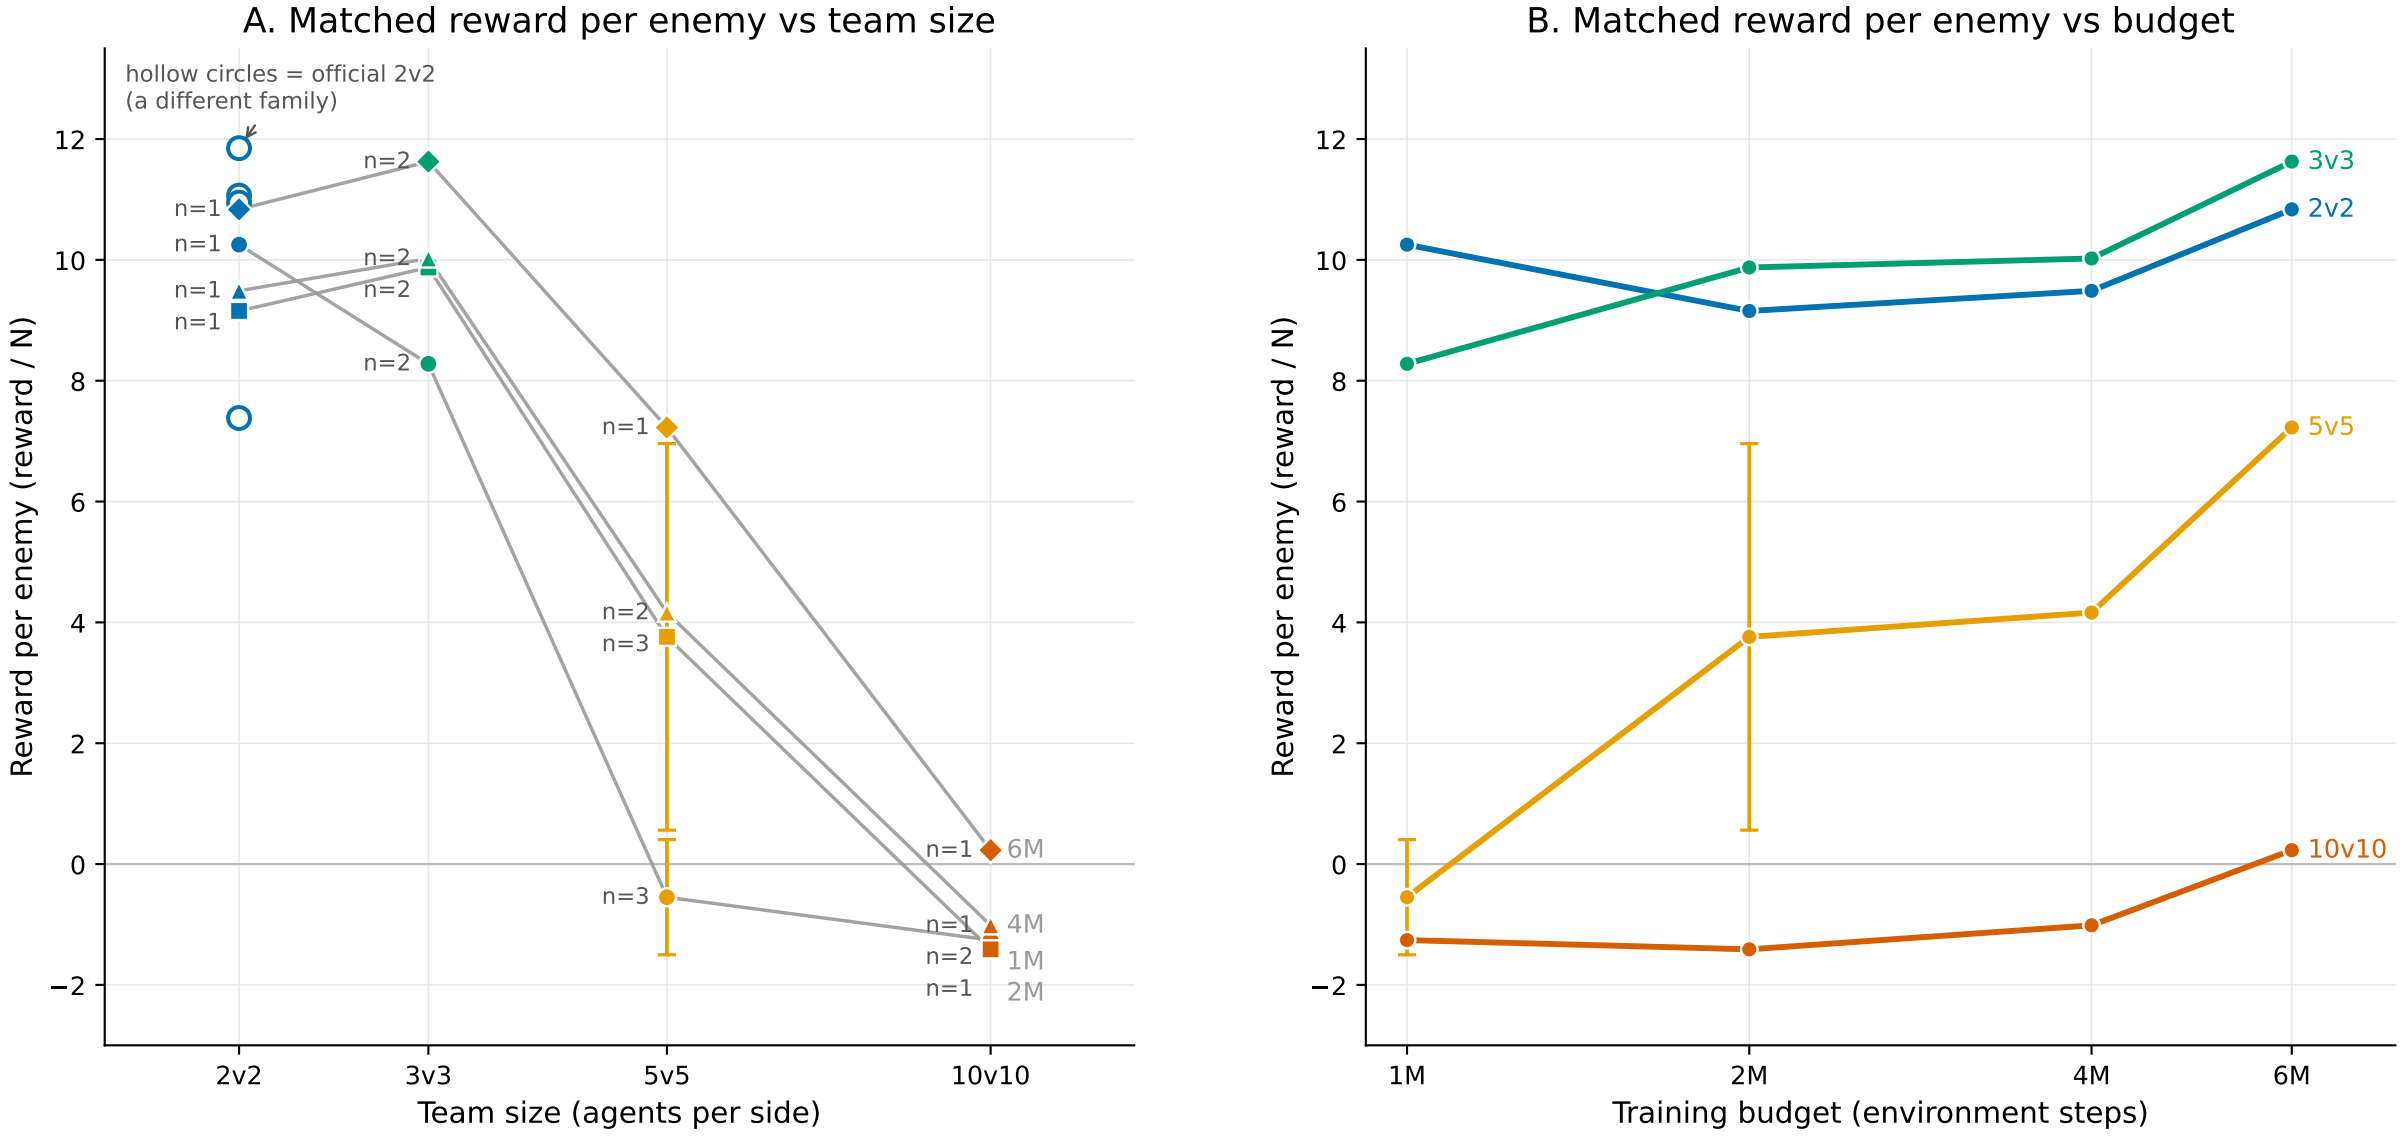

Matched value = mean reward per enemy over the fixed [budget - 200k, budget] env-step bracket (data.MATCHED_WINDOW) of metrics/normalized_scaling.csv; a budget a run never reached is omitted - nothing is interpolated.
Seeds per point at 1M/2M/4M/6M: 2v2 1/1/1/1, 3v3 2/2/2/2, 5v5 3/3/2/1, 10v10 2/1/1/1. Error bars = +/-1 sample s.d. across seeds, drawn only where >= 3 seeds exist; points on 1-2 seeds carry no interval.
Panel A: shape = budget (o 1M, s 2M, ^ 4M, D 6M), colour = team size, grey line = the budget's trajectory across team sizes. Hollow circles = the official-config 2v2 run (a different config family, 1 seed) at the same four budgets, shown for reference only; the official 1v1 / SingleControl runs are excluded (different tasks, reward scales an order of magnitude apart).
LIMITS: reward/N normalisation is only meaningful inside the NoWeapon family (its PostureReward accumulates over enemy aircraft) and is not a unified task-success metric; different N are different tasks (scale, reward make-up and difficulty all differ), so more reward per enemy does not mean a higher win rate. No reliable win-rate or common evaluation metric exists for these runs - this is an exploratory comparison of training reward.
Lines join the means over the seeds available at each budget, so a point's seed set can shrink at later budgets; no run or seed is dropped or filled in.